

Project 03

Virtual Memory & File system

Due date
2025. 06. 18 23:59

Overview

- Copy-on-Write Fork
 - Memory management optimization through lazy allocation
- Large Files
 - File system extension with doubly-indirect blocks
- Symbolic Links
 - Advanced file referencing and path resolution

Project 03

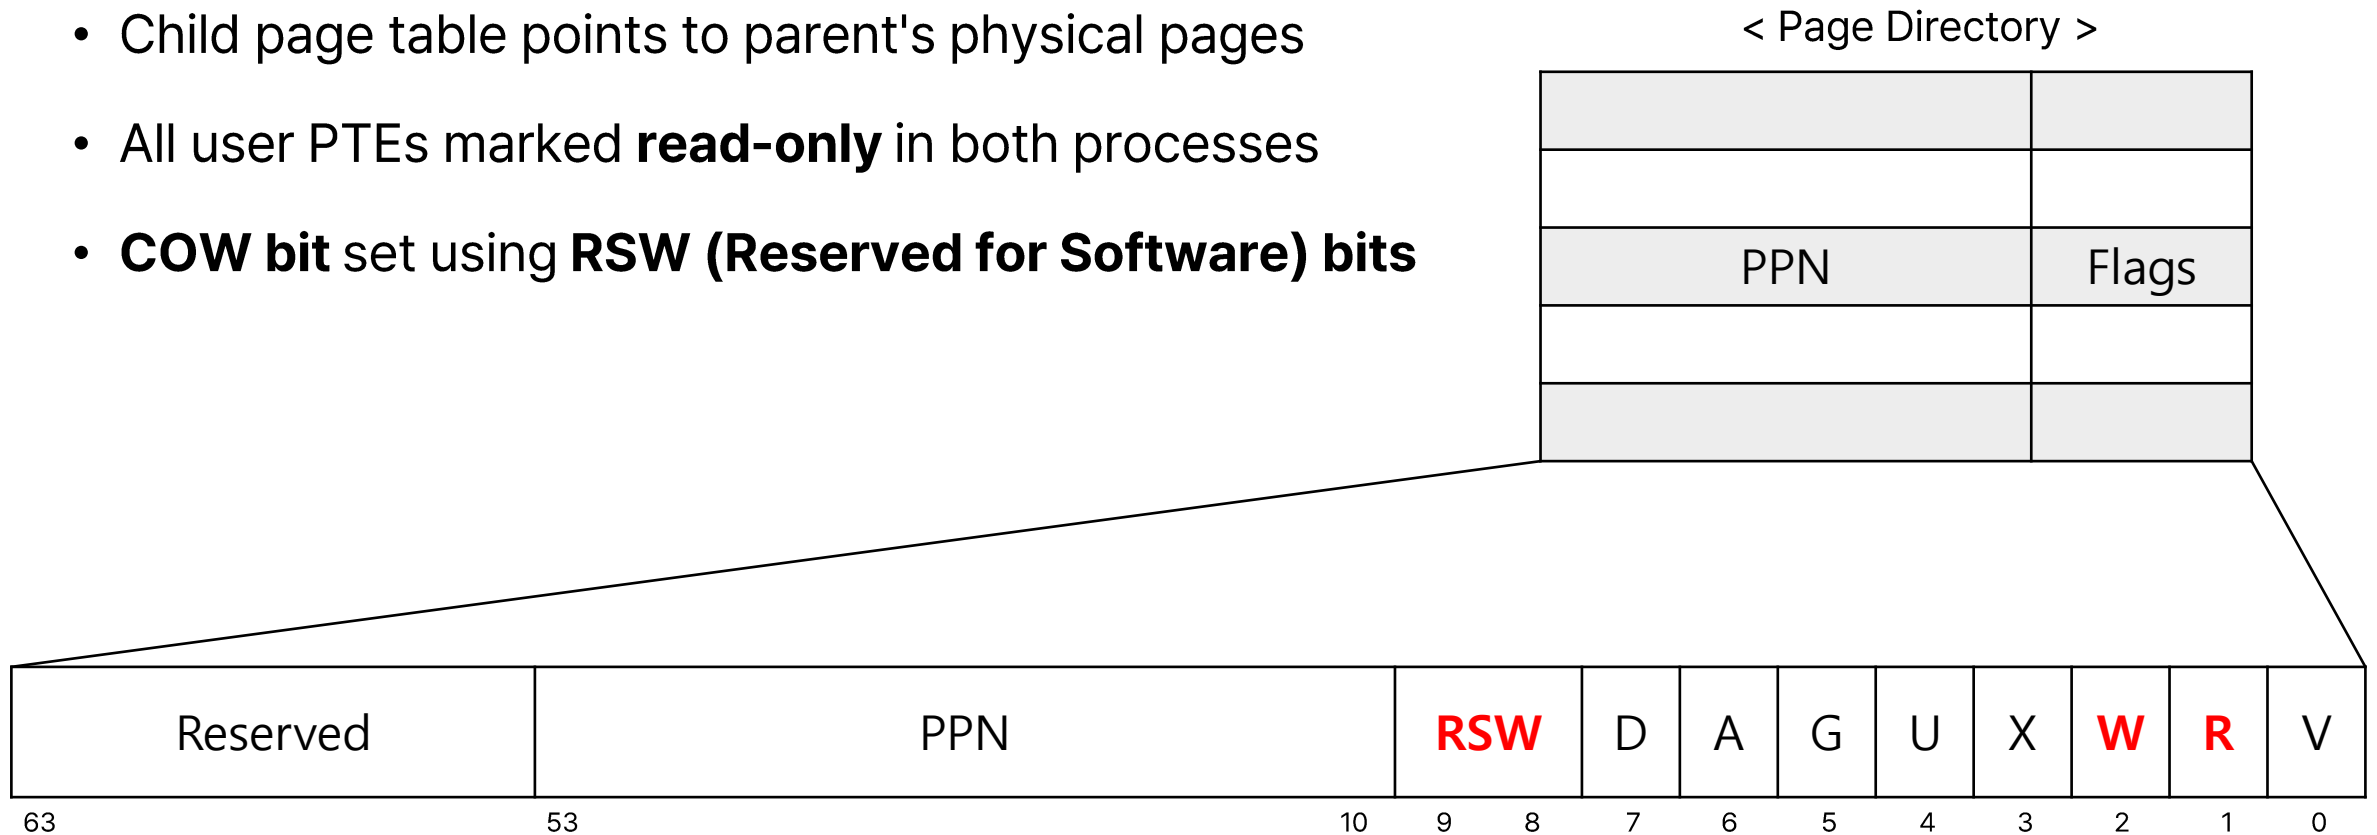
Copy-on-Write Fork

Overview

- Problem: Inefficient memory copying in fork()
- Current fork() copies all parent memory immediately
 - Wasteful when followed by exec()
 - Can fail due to insufficient memory
- Solution: Copy-on-Write (COW)
 - Share pages between parent and child initially
 - Copy pages only when modified

COW Fork Architecture

- Initial State:
 - Child page table points to parent's physical pages
 - All user PTEs marked **read-only** in both processes
 - COW bit** set using **RSW (Reserved for Software) bits**



COW Fork Architecture

- Write Access:
 - **Page fault** triggered on write attempt
 - Kernel allocates new physical page
 - Original page copied to new page
 - PTE updated with write permission
- Reference Counting:
 - Track number of processes sharing each page
 - Free pages only when reference count reaches zero

Implementation Tasks

- Core Functions to Modify:
 - `uvmcopy()`: Share pages instead of copying
 - `usertrap()`: Handle COW page faults
 - `copyout()`: Apply COW logic to kernel operations
 - `kalloc()/kfree()`: Implement reference counting
- Key Challenges:
 - Distinguishing COW faults from other page faults
 - Proper synchronization for reference counts
 - Handling out-of-memory conditions
 - Maintaining compatibility with existing code

Cow Test Result

- simpletest() - Basic COW Memory Allocation and Behavior Verification
- threetest() - Multi-Process COW Stress Testing
- filetest() - System Call Integration with COW (copyout()) Compatibility

```
$ cowtest
simple: ok
simple: ok
three: ok
three: ok
three: ok
file: ok
ALL COW TESTS PASSED
$ █
```


Project 03

Large files

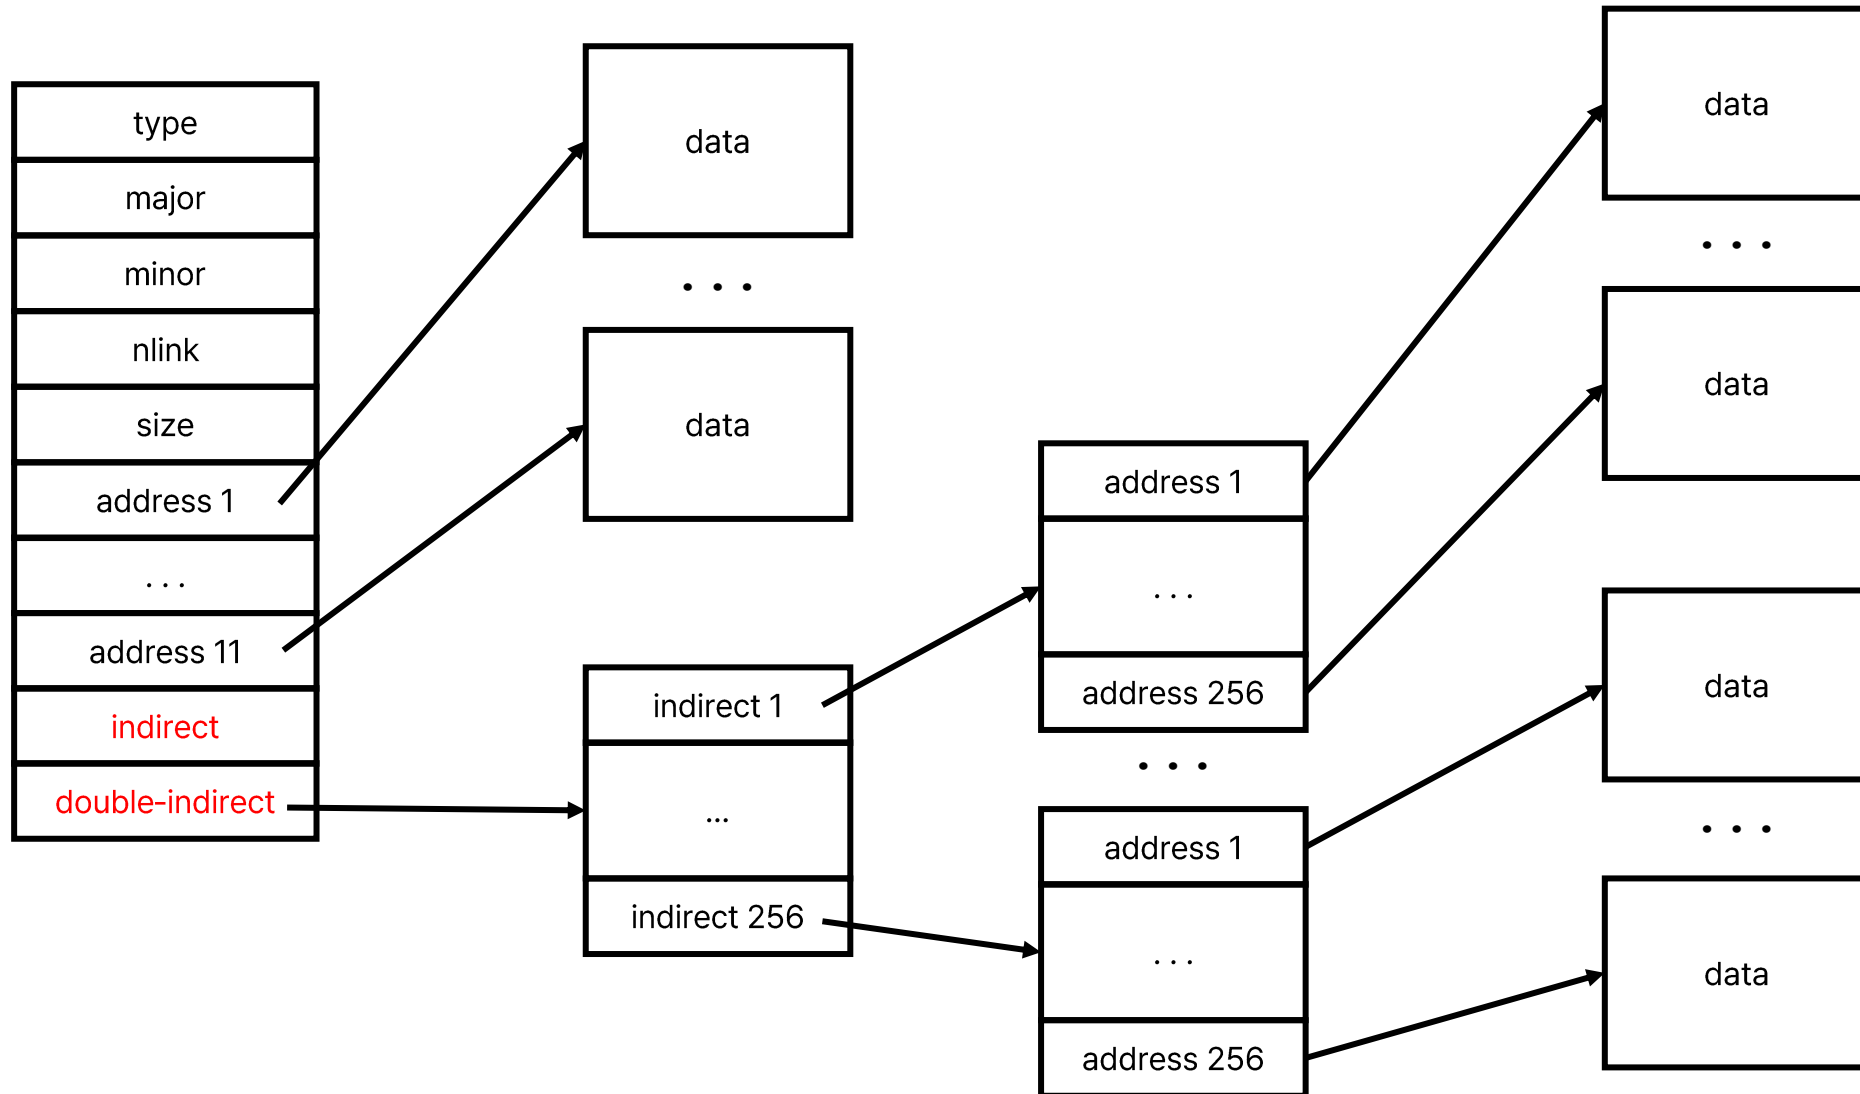
Overview

- Problem: Limited file size in xv6
 - Current limit: 268 blocks (268 KB)
 - Structure: 12 direct + 1 singly-indirect
 - Insufficient for modern applications
- Solution: **Doubly-Indirect Blocks**
 - Support files up to 65,803 blocks (~64 MB)
 - Structure: 11 direct + 1 singly-indirect + 1 doubly-indirect
 - Maintains backward compatibility

Current Block Addressing Hierarchy



Block Addressing Hierarchy



Implementation Tasks

- Constants and Structures:
 - Change FSSIZE from 2000 to 200000
 - Change NDIRECT from 12 to 11
 - Update MAXFILE calculation
 - Modify `addrs[]` array in `dinode` and `inode` structures
- Core Functions:
 - `bmap()`: Implement doubly-indirect address translation
 - `itrunc()`: Free doubly-indirect blocks properly
 - `create()`: Support new `inode` structure Critical

Largefile Test Result

- Large File Creation and Sequential Writing
- File Size Validation
- Data Integrity Verification

```
$ bigfile
```

```
.....  
.....  
.....  
.....  
.....
```

```
wrote 65803 blocks
```

```
bigfile done; ok
```


```
$ █
```

Project 03

Symbolic links

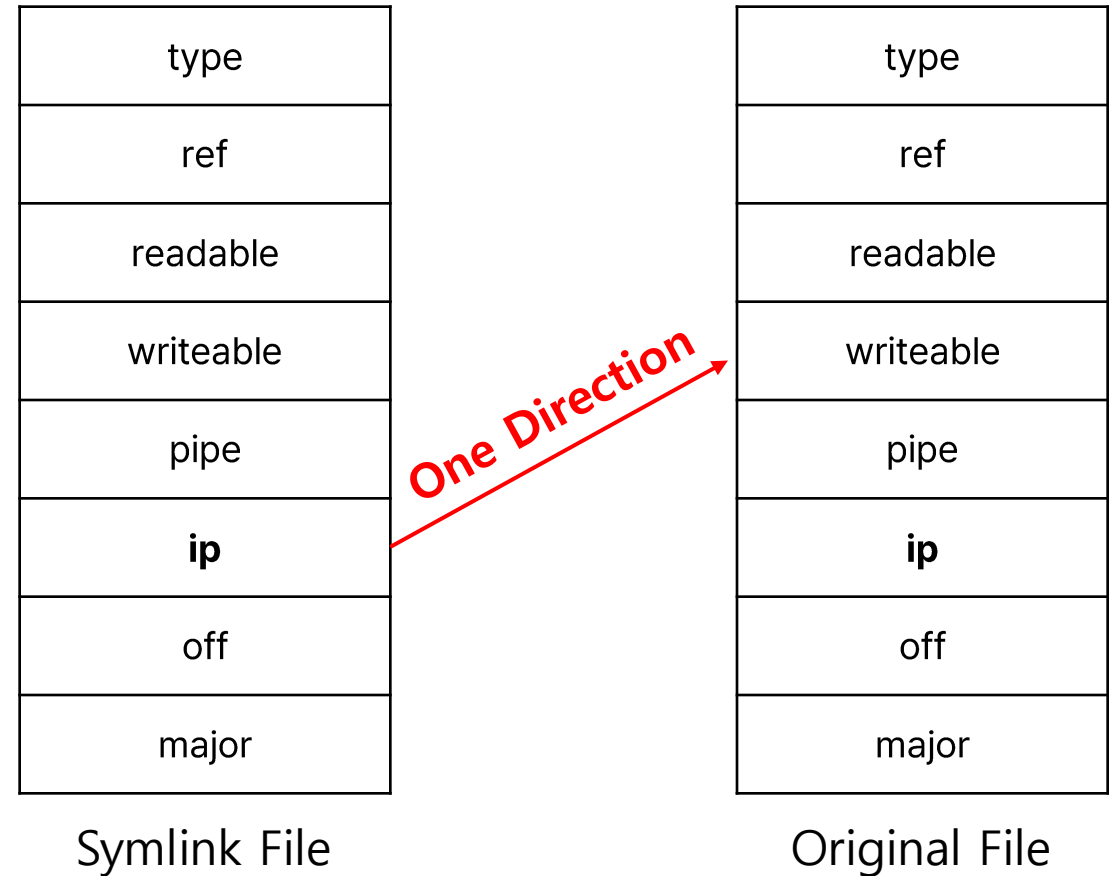
Overview

- Problem: Rigid file referencing in xv6
 - Only hard links available (inode-based)
 - Cannot reference across file systems
 - No support for broken or flexible links



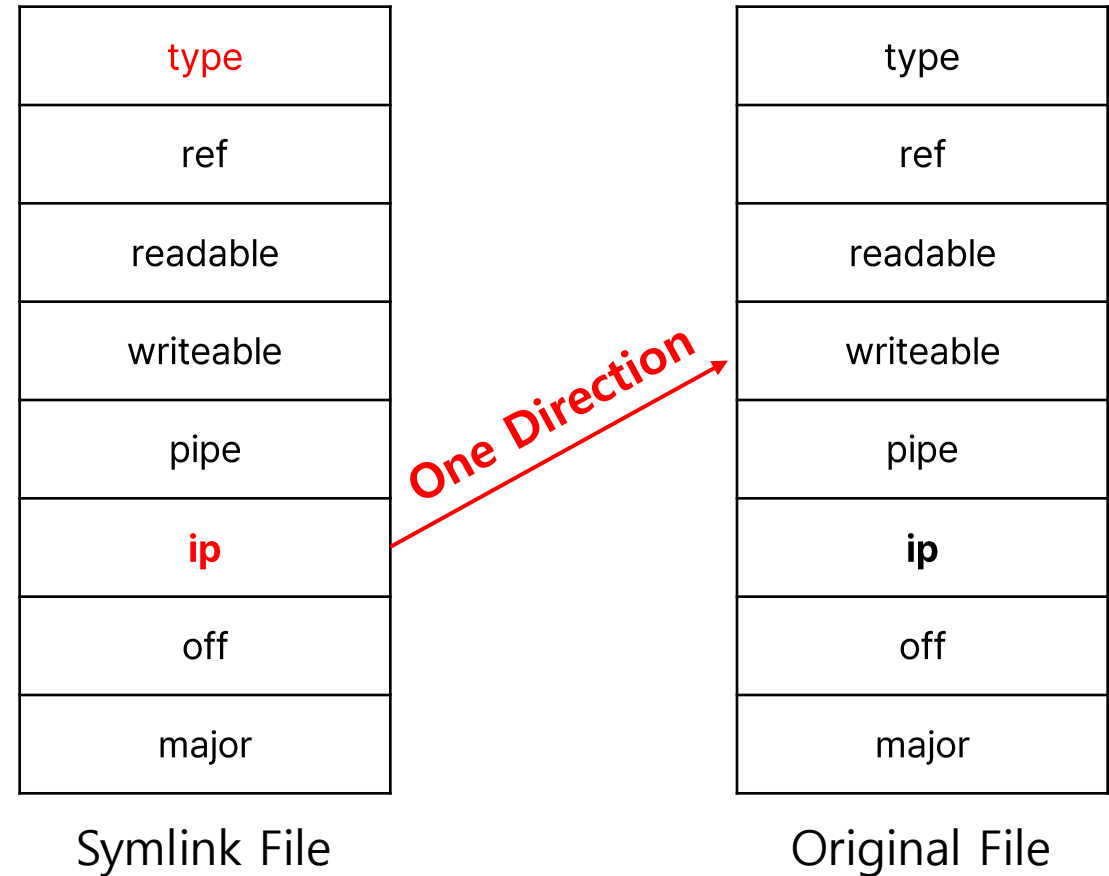
Overview

- Problem: Rigid file referencing in xv6
 - Only hard links available (inode-based)
 - Cannot reference across file systems
 - No support for broken or flexible links
- Solution: **Symbolic Links**
 - Pathname-based file references
 - Can point to non-existent files
 - Support for cross-filesystem references
 - Enhanced file system flexibility



Symbolic Links Architecture

- File Type Extension:
 - New T_SYMLINK file type
 - Target path stored in symlink's data blocks
 - Distinguished from regular files and directories
- System Call Interface:
 - `symlink(target, path)`: Create symbolic link
 - Modified `open()`: Follow or access links directly
 - `O_NOFOLLOW` flag: Open link itself, not target
- Path Resolution:
 - Recursive link following (max depth: 10)
 - Cycle detection to prevent infinite loops
 - Proper error handling for broken links



Implementation Tasks

- New Components:
 - T_SYMLINK file type in kernel/stat.h
 - O_NOFOLLOW flag in kernel/fcntl.h
 - symlink() system call implementation
- Modified Functions:
 - sys_open(): Implement link resolution logic
 - create(): Support symbolic link creation

Symlink Test Result

- testsymlink() - Core Functionality Testing
 - Basic Link Creation and Access
 - Broken Link Handling
 - Circular Reference Detection
 - Non-existent Target Linking
 - Chain Link Resolution
- concur() - Concurrent Testing

```
$ symlinktest
Start: test symlinks
test symlinks: ok
Start: test concurrent symlinks
test concurrent symlinks: ok
$
```

Evaluation

- **Completeness** The xv6 operating system must function correctly according to the specification requirements.
- **Wiki & Comment** Grading will be based on the wiki documentation, so the wiki should be written in as much detail as possible.
- **Deadline** The submission deadline must be strictly observed. After the deadline, your GitHub writing permissions will be revoked.
- **DO NOT SHARE AND COPY!!**

Wiki

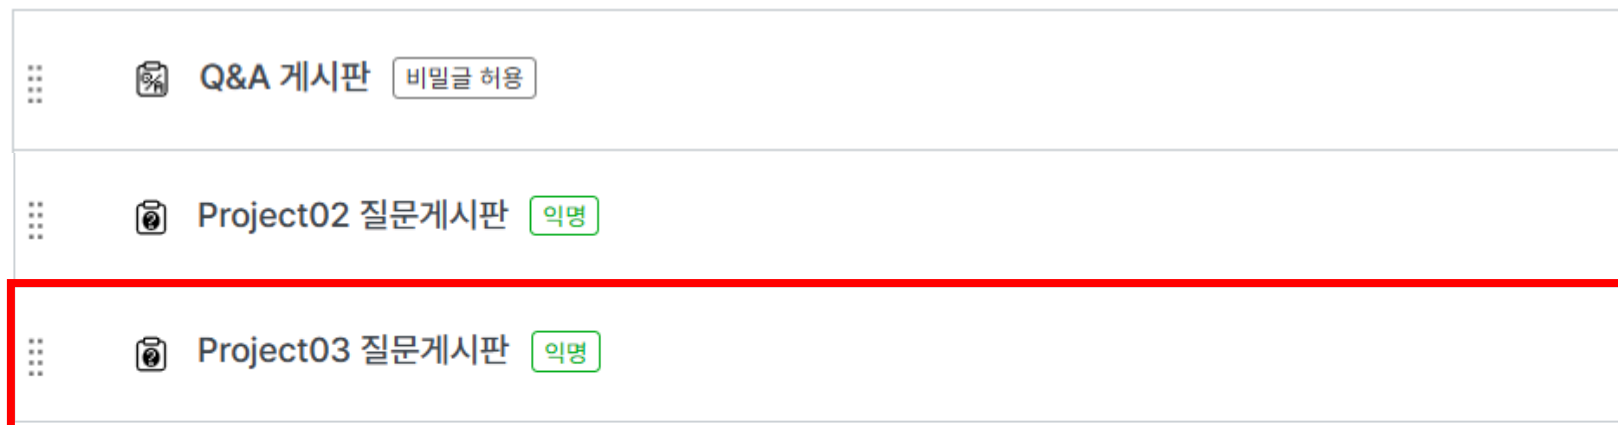
- **Design** Outline your implementation approach for meeting the project requirements
- **Implementation** Explain key code modifications and their purpose, focusing on changes from the original code.
- **Results** Show evidence of successful implementation with:
Compilation process, Screenshots of working code, Explanation of program flow
- **Troubleshooting** Describe any problems encountered, solutions applied, and any unresolved issues.
- Additional content may be included if relevant.

Submission

- Submit your implemented code and wiki through GitHub.
 - **Refer to the announcement and create a new repository.**
 - Rename the repository to "**project03-[student ID]**"
- The wiki file should be named "**OS_project03_[class number]_[student ID].pdf**".
- Submission deadline: **June 18, 2025, 23:59**
 - Late submissions will be accepted via **email** until **June 19, 2025, 23:59**, but will only receive **50%** of the possible score.

Q&A

- For questions related to the project, please use the question board (Project 03 Question Board) on the LMS.
- Questions sent by email will **not** be answered.
- For questions not related to the project, please use the Q&A board or send an email.



Q & A